



Docker Best Practice

Docker basic level II

Why bother?

- Writing a good Dockerfile will result in docker container images, which are:
 - Smaller
 - Cleaner
 - Secure
- To archive these, we write Dockerfiles which follow the best practice, i.e.:
 - They are cleanly written - thus easy to read / debug
 - They are organized in logical sections
 - They have instructions written in logical order
 - ... and more (covered next)


Re-cap: Docker container layers

\$ cat Dockerfile

FROM centos:7

RUN yum -y install emacs

RUN yum -y instal httpd



\$ docker build -t local/centos-apache .


\$ docker image ls

local/centos-apache	latest	c8b3036d9c8a	5 seconds ago	631MB
centos	latest	5e2539a48d9s	5 minutes ago	300MB

Re-cap: Docker container layers ...

```
$ docker image history local/centos-apache
```

IMAGE	CREATED	CREATED BY	SIZE
c8b3036d9c8a	9 minutes ago	/bin/sh -c yum -y install httpd	631MB
22b134c9419c	19 minutes ago	/bin/sh -c yum -y install emacs	350MB
5e2539a48d9s	1 months ago	/bin/sh -c #(noh) CMD ["/bin/bash"]	300MB
<missing>	1 months ago	/bin/sh -c #(noh) LABEL org.label-scheme..	0B
<missing>	1 months ago	/bin/sh -c #(noh) ADD file:14a318104952b..	203B



AUFS - Advanced (multi-layered) Unification FileSystem

Problem:

- You have 3 directories, dir1, dir2 and mountdir.
- You want to somehow mount dir1 and dir2 in mountdir, so you can see (combined) contents of dir1 and dir2 in mountdir.

Older Solutions:

- Use unionfs or aufs

Latest Solution:

- AUFS is not used anymore
- Use OverlayFS

OverlayFS - How it works?

- OverlayFS layers two directories on a single Linux host and presents them as a single directory.
- These directories are called **layers** and the unification process is referred to as a **union mount**.
- OverlayFS refers to the lower directory as **lowerdir** and the upper directory a **upperdir**.
- The unified view is exposed through its own directory called merged.
- The overlay2 driver natively supports up to 128 lower OverlayFS layers, and provides better performance.

OverlayFS - Keep packages to the minimum

- When creating container images, start with smallest possible base image based on your requirements
- Install only enough packages which are needed to run your application. I.e.
 - Don't install an entire operating system inside it
 - Don't download the entire internet inside it
- For debugging purposes, attach “tools” containers to your application containers when necessary

OverlayFS - Use unix's philosophy - De-couple applications

- A container should do only one-thing, and do it well
 - For example: Don't combine Apache and MySQL in one container: they should be two separate containers
- Individual containers serving/providing individual services are easier to: manage, scale, troubleshoot, update, upgrade, etc.

Best Practice: Minimize number of layers in an image

- Remember, only the instructions RUN, COPY, ADD create layers
- Consolidate multiple RUN commands into one, or have lesser RUN commands
- Organize files in a better way before COPY-ing into the container image.
- Use COPY instead of ADD, as COPY is more transparent, and has predictable behavior
- Don't use ADD to fetch packages from remote URLs. Use curl or wget instead.

Best Practice: Minimize number of layers in an image

\$ cat Dockerfile.long

```
FROM debian
RUN apt-get update
RUN apt-get install -y vim-common
RUN apt-get install -y curl
RUN apt-get install -y netcat
RUN apt-get install -y tcpdump
RUN apt-get install -y dnsutils
RUN apt-get install -y net-tools
```

\$ cat Dockerfile.compact

```
FROM debian
RUN apt-get update \
    && apt-get -y install vim-common \
    curl netcat tcpdump dnsutils \
    net-tools
```

```
$ docker build -f Dockerfile.long -t local/simple:long . --no-cache
$ docker build -f Dockerfile.compact -t local/simple:compact . --no-cache
```

\$ docker image ls

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
local/simple	compact	372v7c6789sb8	1 minutes ago	229MB
local/simple	simple	876252g2u752b	2 minutes ago	239MB

Best practice: Use multi-stage builds

- Only copy the artifacts you need into the final image.

\$ cat Dockerfile

```
FROM golang:1.14 as builder
```

```
WORKDIR /tmp
```


```
COPY myapp.go /tmp
```

```
RUN go build myapp.go
```

```
FROM scratch as myapp
```

```
COPY --from=builder /tmp/myapp /myapp
```

```
CMD ["/myapp"]
```



\$ docker build -t local/myapp . --no-cache

\$ docker image ls | grep golang

golang	1.14	12a265b042a	9 days ago	809MB
--------	------	-------------	------------	-------

\$ docker image ls | grep myapp

local/myapp	latest	fbf7826bs09c	20 seconds ago	2.07MB
-------------	--------	--------------	----------------	--------

Best practice: sort multi-line arguments

- For better readability, prevent double entries and debugging sort all command arguments alphabetically.

\$ cat Dockerfile

```
FROM      centos:7
RUN       yum -y install \
          vim-common \
          curl \
          socat \
          traceroute \
          net-tools \
          bind-utils \
          tcpdump
```

\$ cat Dockerfile

```
FROM      centos:7
RUN       yum -y install \
          bind-utils \
          curl \
          net-tools \
          socat \
          tcpdump
          traceroute \
          vim-common
```

Best practice: “Debian” & derivatives - Avoid RUN apt-get upgrade and dist-upgrade

- Many of these “essential” packages from the parent images cannot upgrade inside an unprivileged container.
- If a package contained in the parent image is out-of-date, contact its maintainers.
- Use `apt-get install -y <package>` to update any particular package automatically.

\$ cat Dockerfile

```
FROM      debian
RUN      apt-get update \
&& apt-get upgrade \
&& apt-get dist-upgrade \
&& apt-get -y install curl \
&& dns-utils \
        net-tools \
        netcat \
```

\$ cat Dockerfile

```
FROM      debian
RUN      apt-get update \
&& apt-get -y install curl \
        dns-utils \
        netcat \
        net-tools \
```

Best practice: “debian & derivatives - Always have “apt-get update” and “apt-get install” in one RUN command

- Using apt-get update in a separate RUN statement causes caching issues
- If apt-get update is a separate command, it pulls in apt package index, and docker stores it as a layer.
- For every package in need of installation, apt-get consults this package index. It may be “now”, or “later in time”.
- When “later in time” comes, the package index - which is cached layer - is old, which can cause an older version of the package to be installed!

<pre>\$ cat Dockerfile FROM debian RUN apt-get update RUN apt-get install -y curl</pre>	<pre>\$ cat Dockerfile RUN apt-get update \ && apt-get install -y curl nginx</pre>
<pre>\$ cat Dockerfile RUN apt-get update RUN apt-get install -y curl nginx</pre>	

Best practice : Use CMD, Entrypoint correctly

- CMD should almost always be used in the form of CMD ["executable","param1","param2" ...]
 - E.g. if the image is for a service, such as Apache, CMD should be ["apache2","-DFOREGROUND"]
- Entrypoint should only be used to "prepare" the container, not run some command as default.
 - ...unless that is the sole purpose of that container, such as a container used as an executable.

```
$ cat Dockerfile
```

```
FROM httpd:2.4
```

```
...
```

```
CMD ["apache2", "-DFOREGROUND"]
```

```
$ cat Dockerfile
```

```
...
```

```
COPY entrypoint.sh /
```

```
ENTRYPOINT ["/entrypoint.sh"]
```

```
CMD ["apache2", "-DFOREGROUND"]
```

```
$ cat Dockerfile
```

```
#!/bin/bash
```

```
Echo "This is a web host : $(hostname)" >> /var/www/html/index.html
```

```
Chmod +r
```

Best practice: be careful in using pipe(|) in RUN commands

- Docker executes these commands the `/bin/sh -c`` interpreter, which only evaluates the exit code of the last operation in the pipe to determine success.
 - There is no guarantee if the RUN command shown below will always work as expected.

\$ cat Dockerfile

FROM busybox

...

RUN wget -O - https://somerandom.site | wc -l > /number

...

- If you want the command to fail due to an error at any stage in the pipe, prepend `'set -o pipefail &&'` to the RUN command.

\$ cat docker-entrypoint.sh

FROM busybox

...

RUN **set -o pipefail** && wget -O - https://somerandom.site | w c -l | /number

Best practice: Avoid running as root

- Running a container as a root does not automatically mean that there is a security problem.
- However, if a service can run as a normal user, then it is best to set a non-root user in the USER directive to ensure that the container will not run as root.
- Avoiding installing or using sudo, as it has unpredictable TTY and signal-forwarding behaviour that can cause problems.
- If you absolutely need functionality similar to sudo, such as initializing the daemon as root but running it as non-root, consider using “gosu”
- To reduce layers and complexity, avoid switching USER back and forth frequently in your Dockerfile.

Best practice: use WORKDIR properly

- For clarity and reliability, you should always use absolute path for your WORKDIR
- You should use WORKDIR instead of proliferating instructions like: ``RUN cd ... && do-something``

\$ cat Dockerfile

```
FROM httpd:2.4
...
COPY web-content.tar /var/www/html
RUN cd /var/www/html
RUN untar web-content.tar
RUN cd ../../..
...
```

\$ cat Dockerfile

```
FROM httpd:2.4
...
WORKDIR /var/www/html
COPY web-content.tar /var/www/html
RUN untar web-content.tar
...
```